

ASSIGNMENT-3

Name: T.Saarvari

Reg.No: 20BCD7057

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Load the dataset

```
In [2]: df = pd.read_csv("housing.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditio
0	13300000	7420	4	2	3	yes	no	no	no	
1	12250000	8960	4	4	4	yes	no	no	no	
2	12250000	9960	3	2	2	yes	no	yes	no	
3	12215000	7500	4	2	2	yes	no	yes	no	
4	11410000	7420	4	1	2	yes	yes	yes	no	

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
----  -
0   price                 545 non-null   int64
1   area                  545 non-null   int64
2   bedrooms              545 non-null   int64
3   bathrooms             545 non-null   int64
4   stories               545 non-null   int64
5   mainroad              545 non-null   object
6   guestroom            545 non-null   object
7   basement              545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning       545 non-null   object
10  parking               545 non-null   int64
11  furnishingstatus      545 non-null   object
dtypes: int64(6), object(6)
memory usage: 51.2+ KB
```

Perform Below Visualizations.

Univariate Analysis

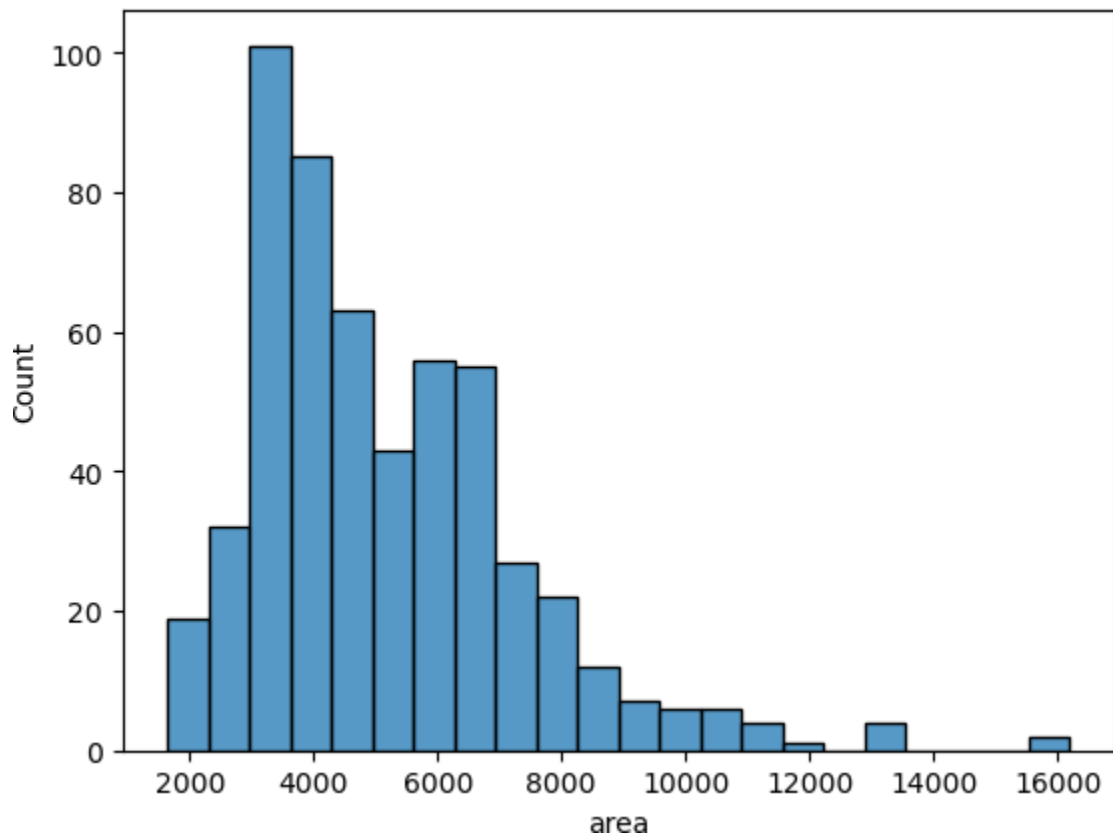
Bi - Variate Analysis

Multi - Variate Analysis

univariate analysis

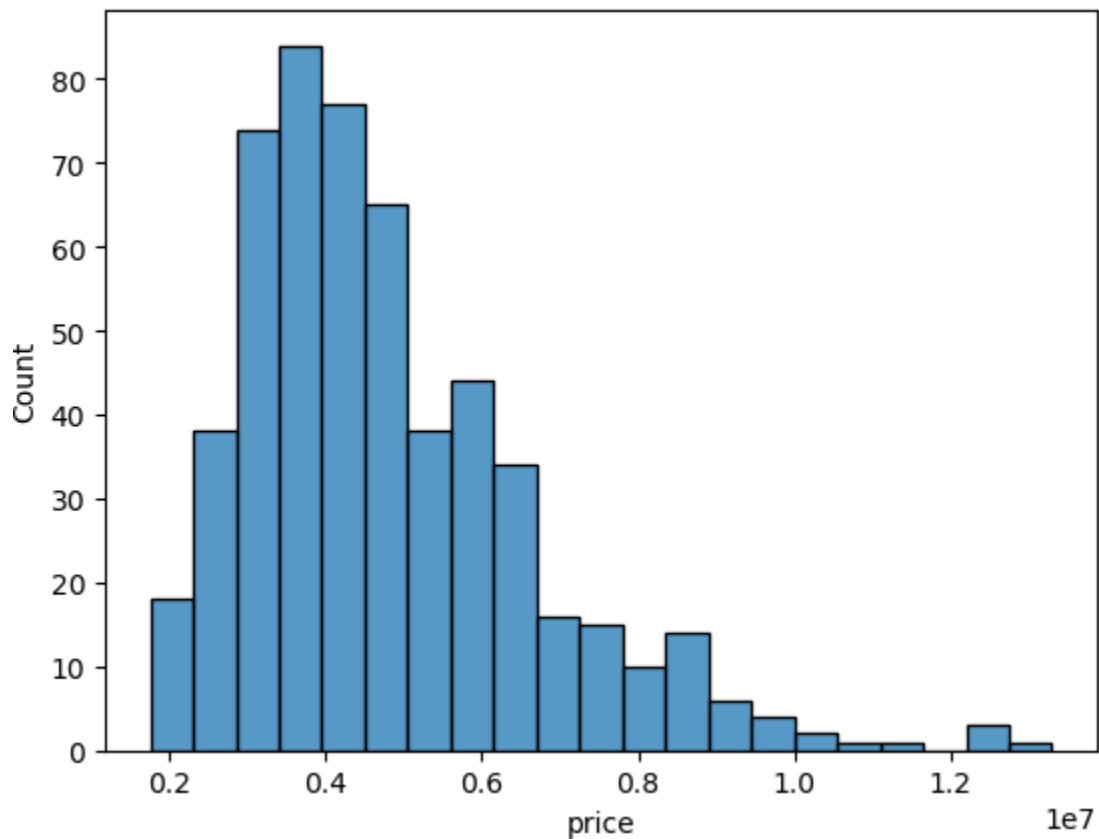
```
In [5]: sns.histplot(df['area'])
```

Out[5]: <Axes: xlabel='area', ylabel='Count'>



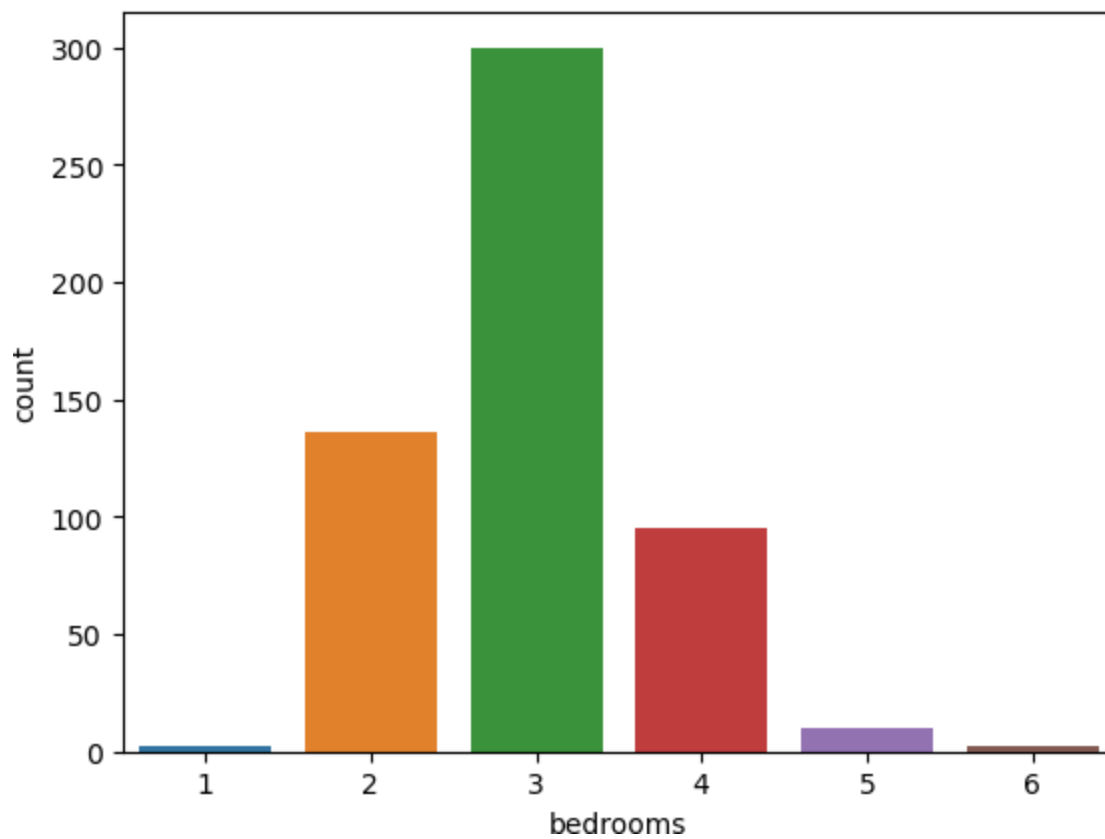
In [6]: `sns.histplot(df['price'])`

Out[6]: <Axes: xlabel='price', ylabel='Count'>



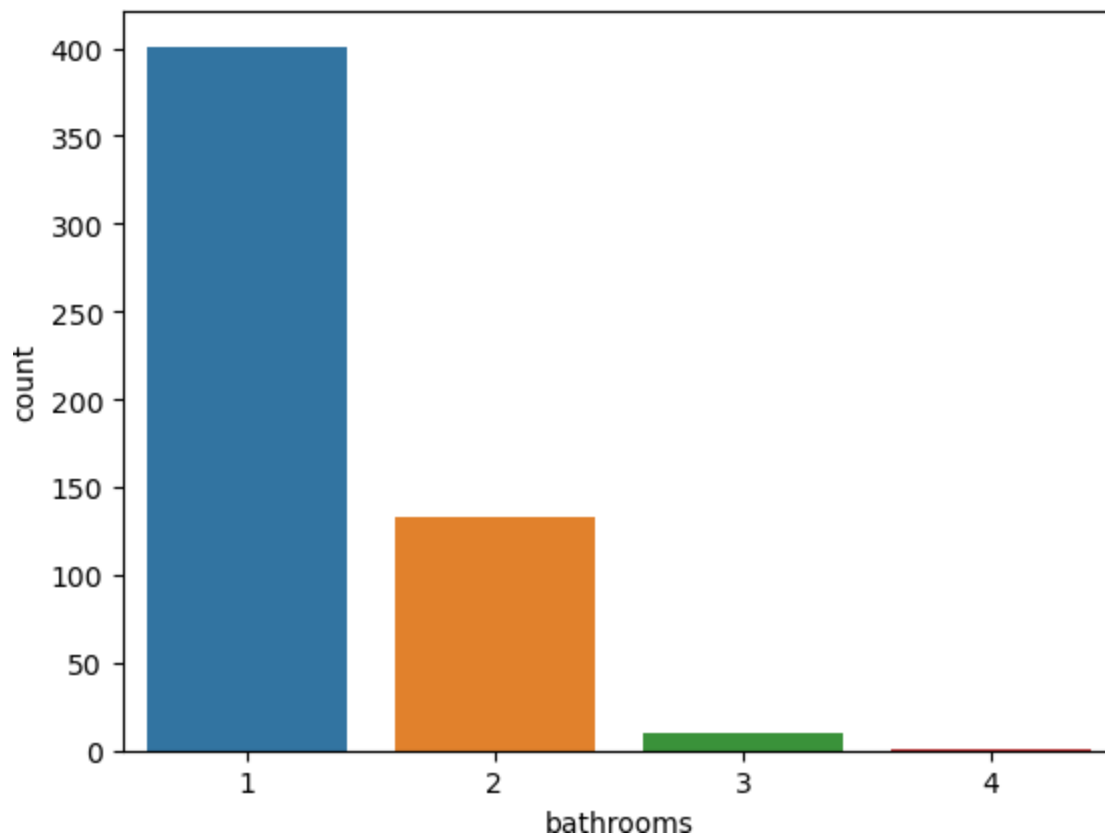
In [7]: `sns.countplot(x = df['bedrooms'])`

Out[7]: <Axes: xlabel='bedrooms', ylabel='count'>



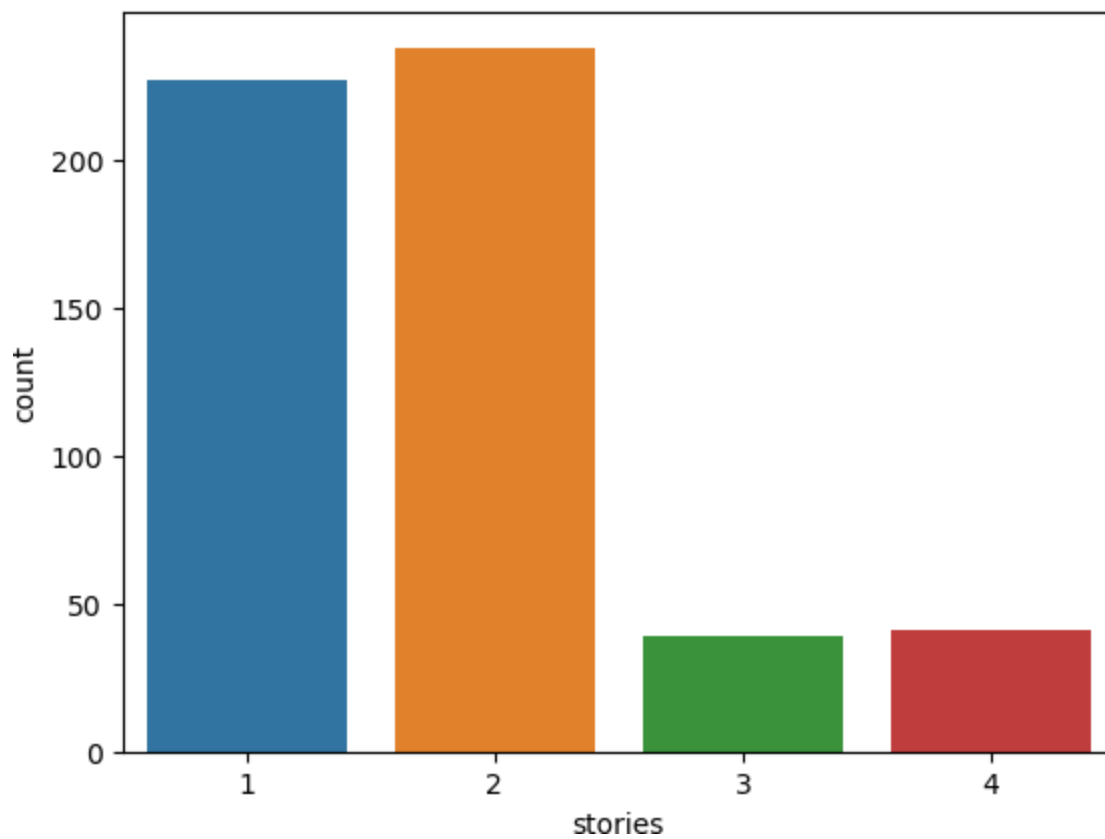
```
In [8]: sns.countplot(x = df['bathrooms'])
```

```
Out[8]: <Axes: xlabel='bathrooms', ylabel='count'>
```



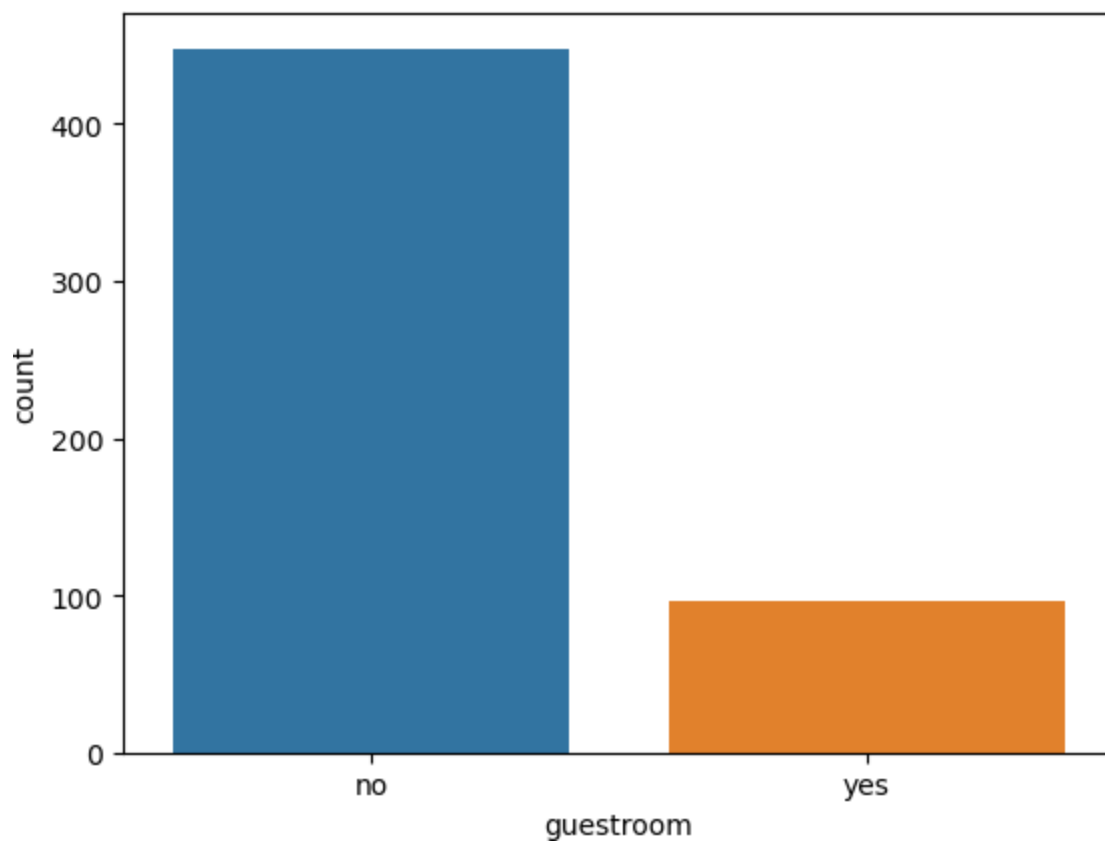
```
In [9]: sns.countplot(x = df['stories'])
```

```
Out[9]: <Axes: xlabel='stories', ylabel='count'>
```



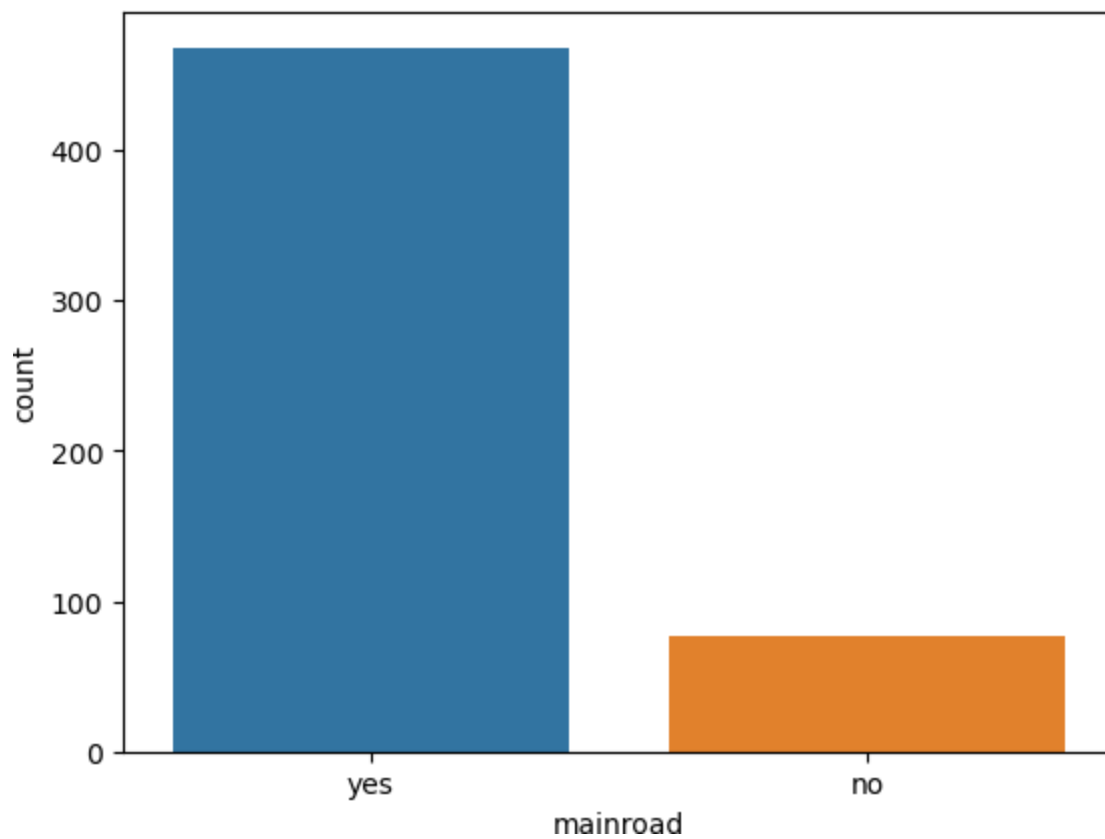
```
In [10]: sns.countplot(x = df['guestroom'])
```

```
Out[10]: <Axes: xlabel='guestroom', ylabel='count'>
```



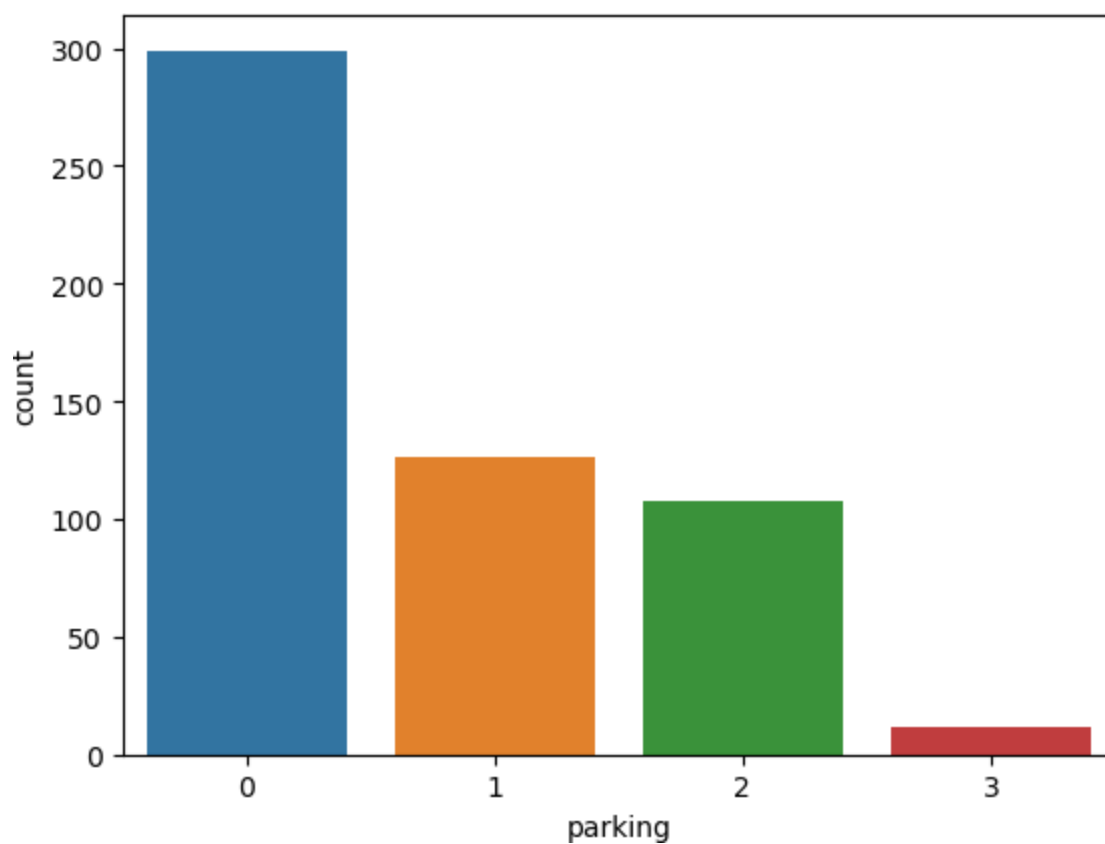
```
In [11]: sns.countplot(x = df['mainroad'])
```

```
Out[11]: <Axes: xlabel='mainroad', ylabel='count'>
```



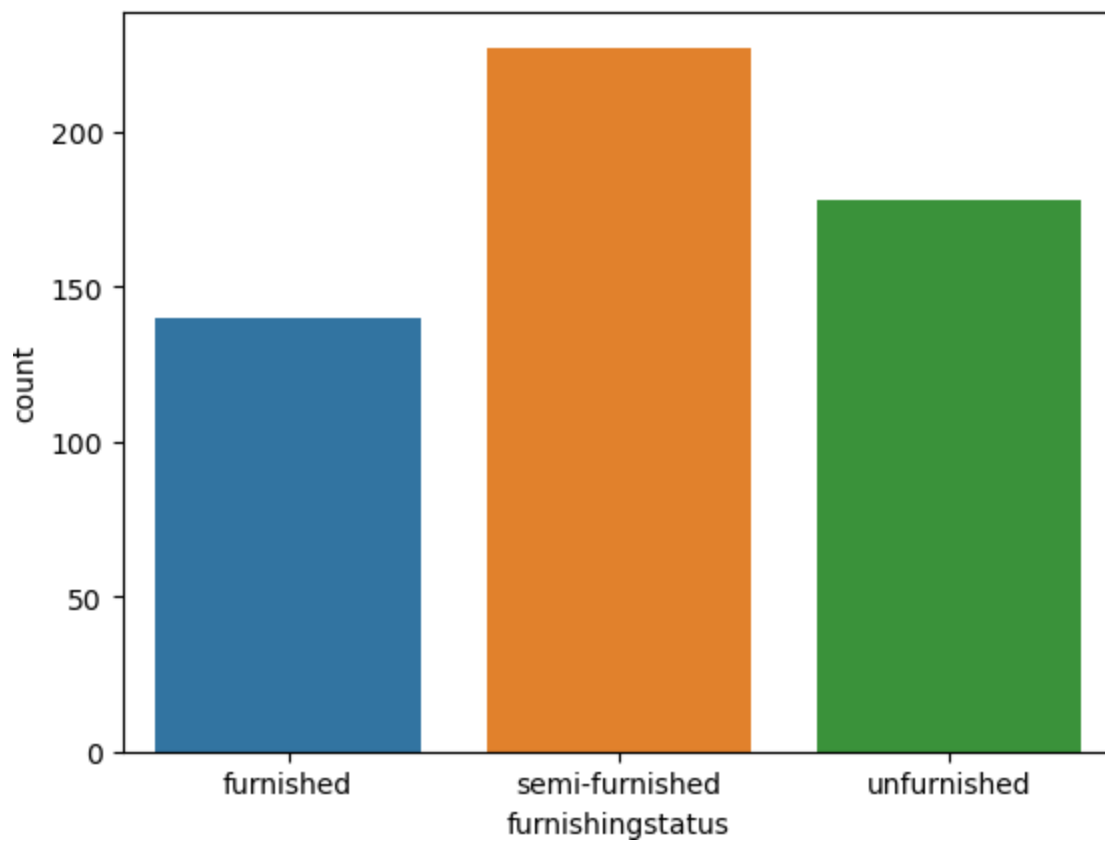
```
In [12]: sns.countplot(x = df['parking'])
```

```
Out[12]: <Axes: xlabel='parking', ylabel='count'>
```



```
In [13]: sns.countplot(x = df['furnishingstatus'])
```

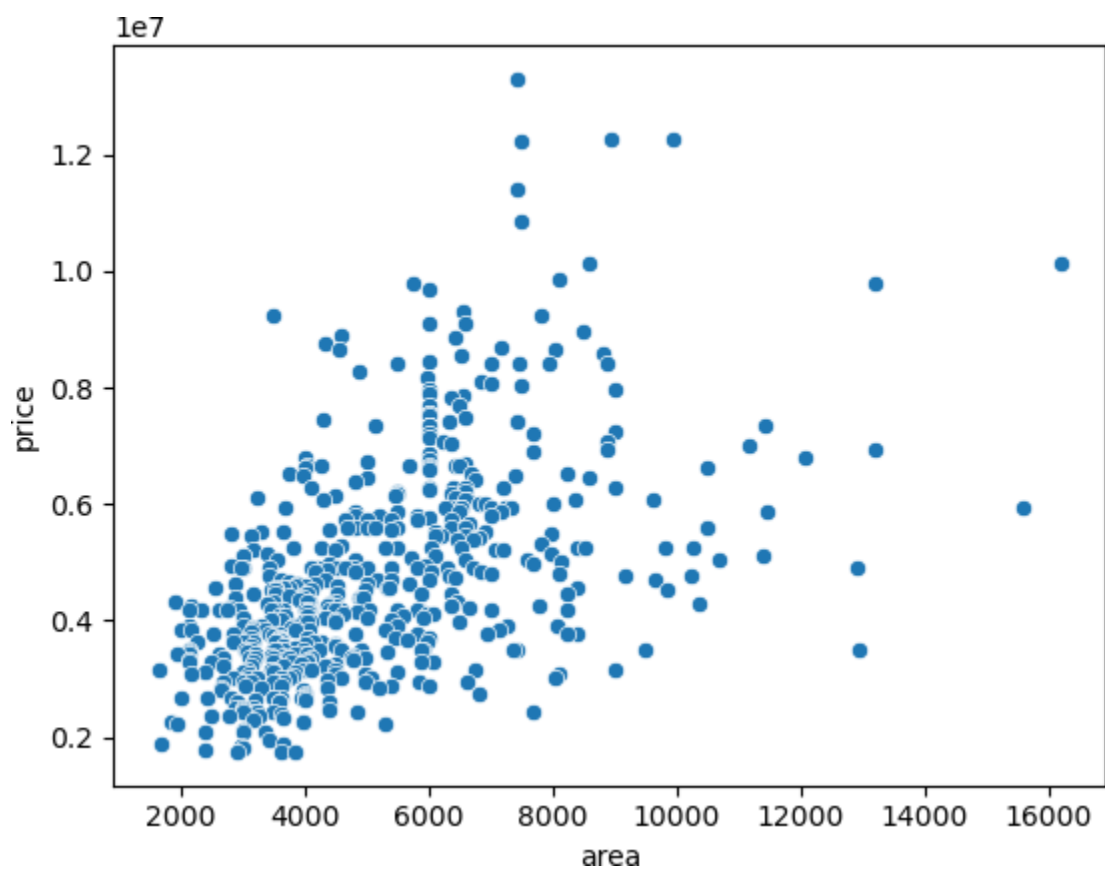
```
Out[13]: <Axes: xlabel='furnishingstatus', ylabel='count'>
```



bivariate analysis


```
In [14]: sns.scatterplot(data = df, x = 'area', y = 'price')
```

```
Out[14]: <Axes: xlabel='area', ylabel='price'>
```




```
In [15]: sns.barplot(data = df, x = 'furnishingstatus', y = 'bedrooms')
```

```
Out[15]: <Axes: xlabel='furnishingstatus', ylabel='bedrooms'>
```



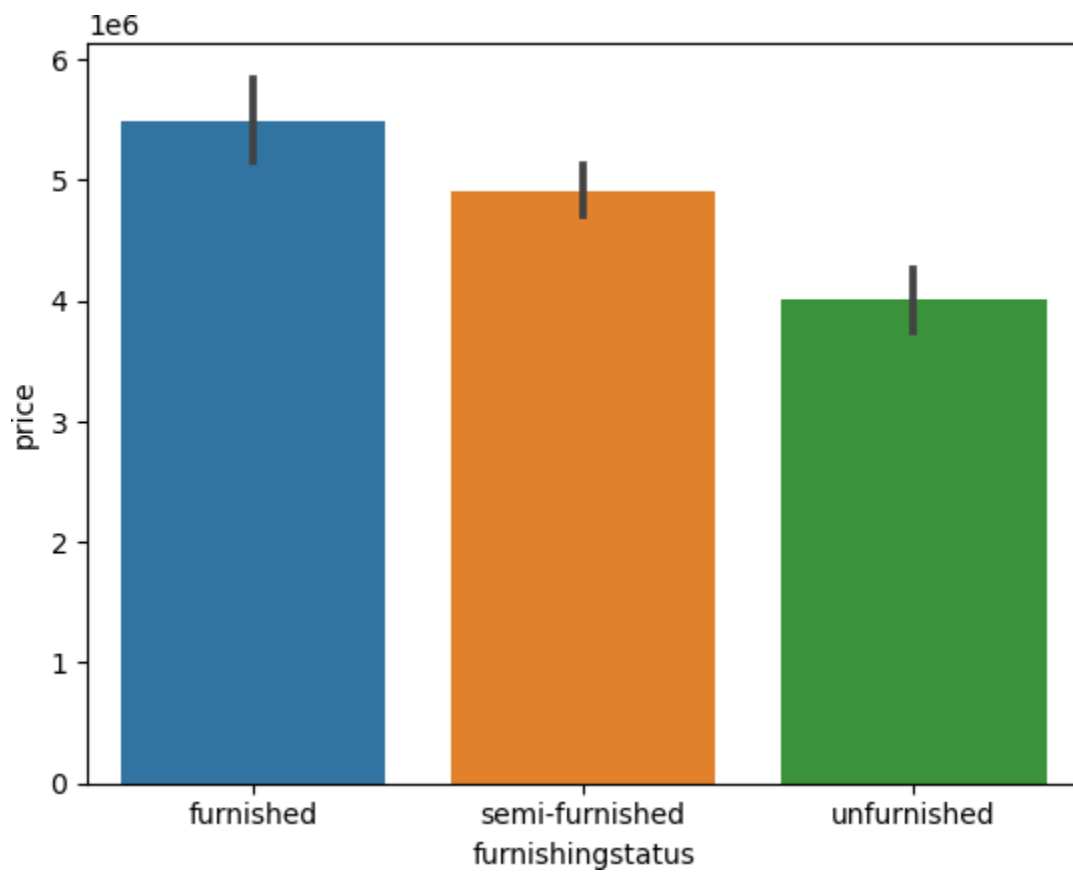
```
In [16]: sns.barplot(data = df, x = 'furnishingstatus', y = 'area')
```

```
Out[16]: <Axes: xlabel='furnishingstatus', ylabel='area'>
```



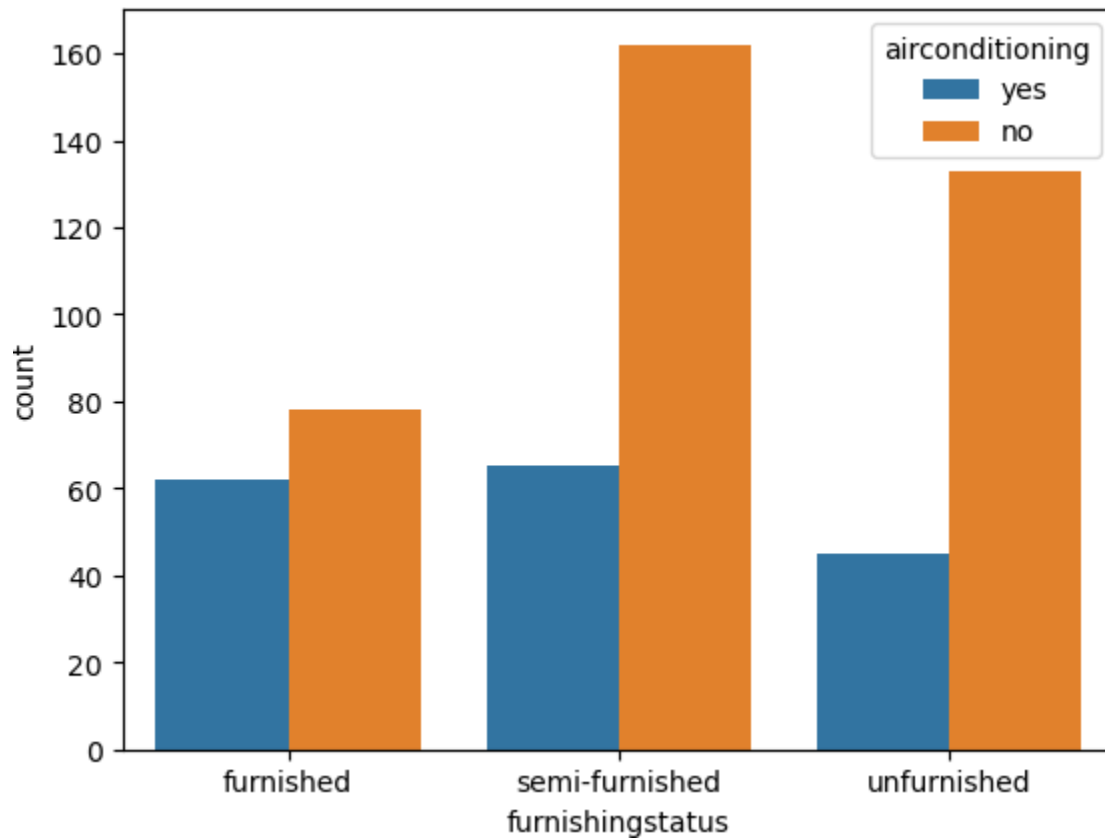
```
In [17]: sns.barplot(data = df, x = 'furnishingstatus', y = 'price')
```

```
Out[17]: <Axes: xlabel='furnishingstatus', ylabel='price'>
```

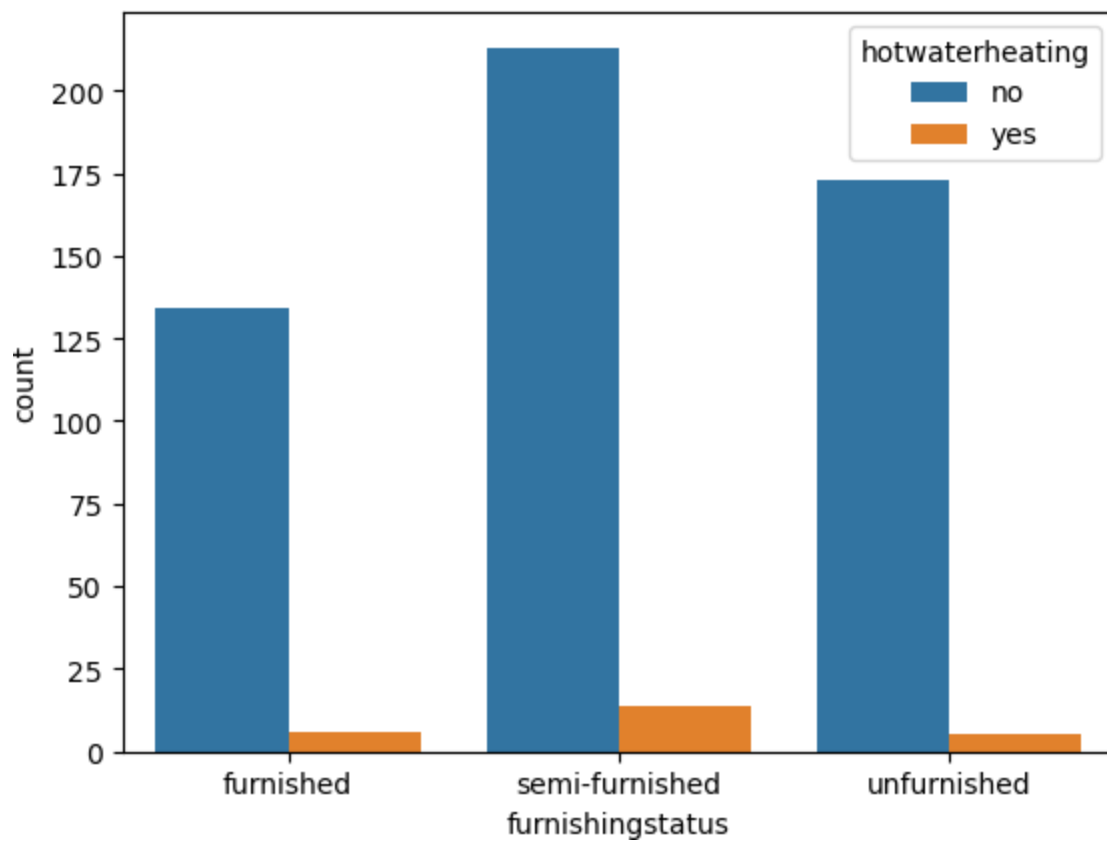
```
In [18]: sns.countplot(x = df['furnishingstatus'], hue = df['airconditioning'])
```

```
Out[18]: <Axes: xlabel='furnishingstatus', ylabel='count'>
```



```
In [19]: sns.countplot(x = df['furnishingstatus'], hue = df['hotwaterheating'])
```

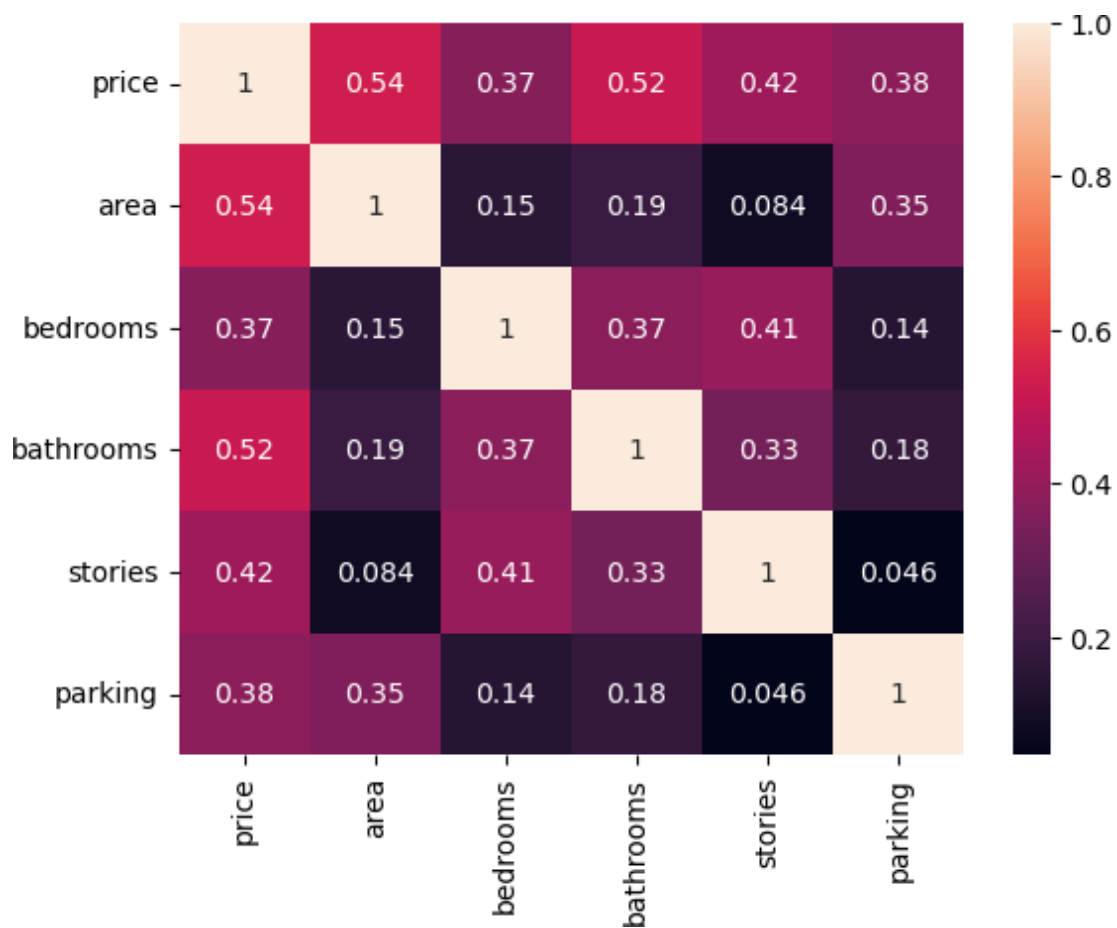
```
Out[19]: <Axes: xlabel='furnishingstatus', ylabel='count'>
```



multivariate analysis

In [20]: `sns.heatmap(df.corr(numeric_only=True), annot = True)`

Out[20]: <Axes: >



Perform descriptive statistics on the dataset.

```
In [21]: df.describe()
```

Out[21]:

	price	area	bedrooms	bathrooms	stories	parking
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

Handle the Missing values.

```
In [22]: df.isnull().sum()
```

Out[22]:

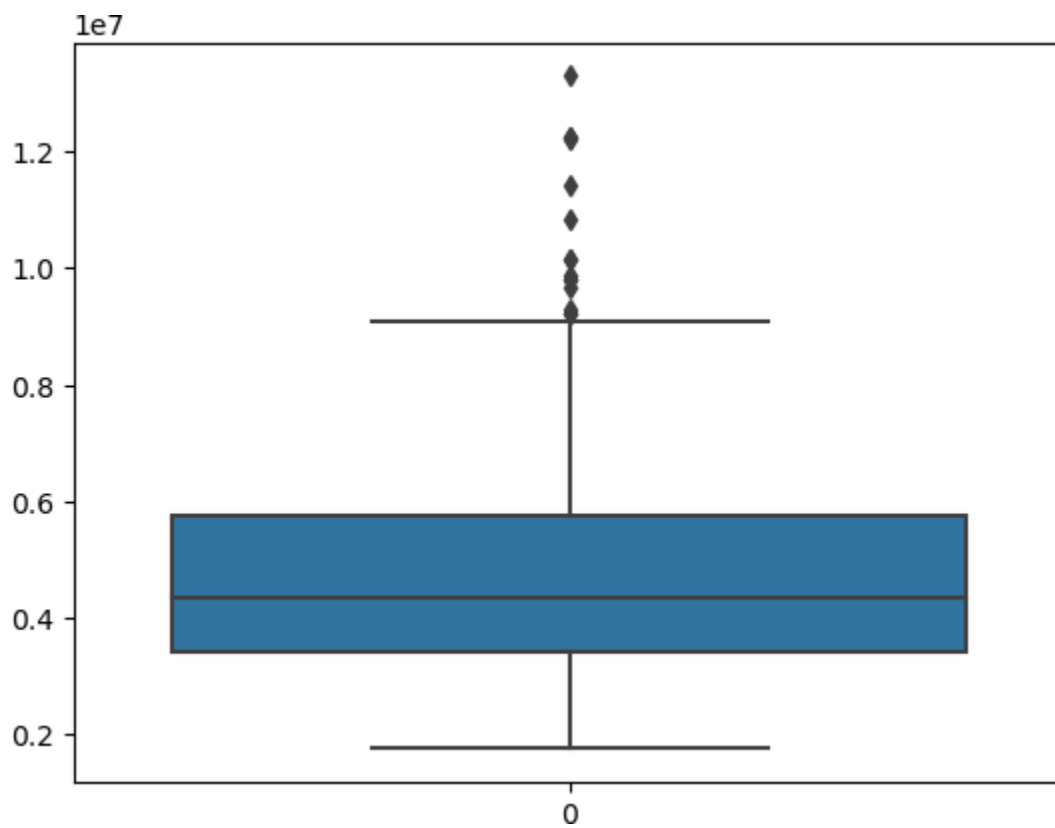
price	0
area	0
bedrooms	0
bathrooms	0
stories	0
mainroad	0
guestroom	0
basement	0
hotwaterheating	0
airconditioning	0
parking	0
furnishingstatus	0
dtype:	int64

Find the outliers and replace the outliers

```
sns.boxplot(df['price'])
```

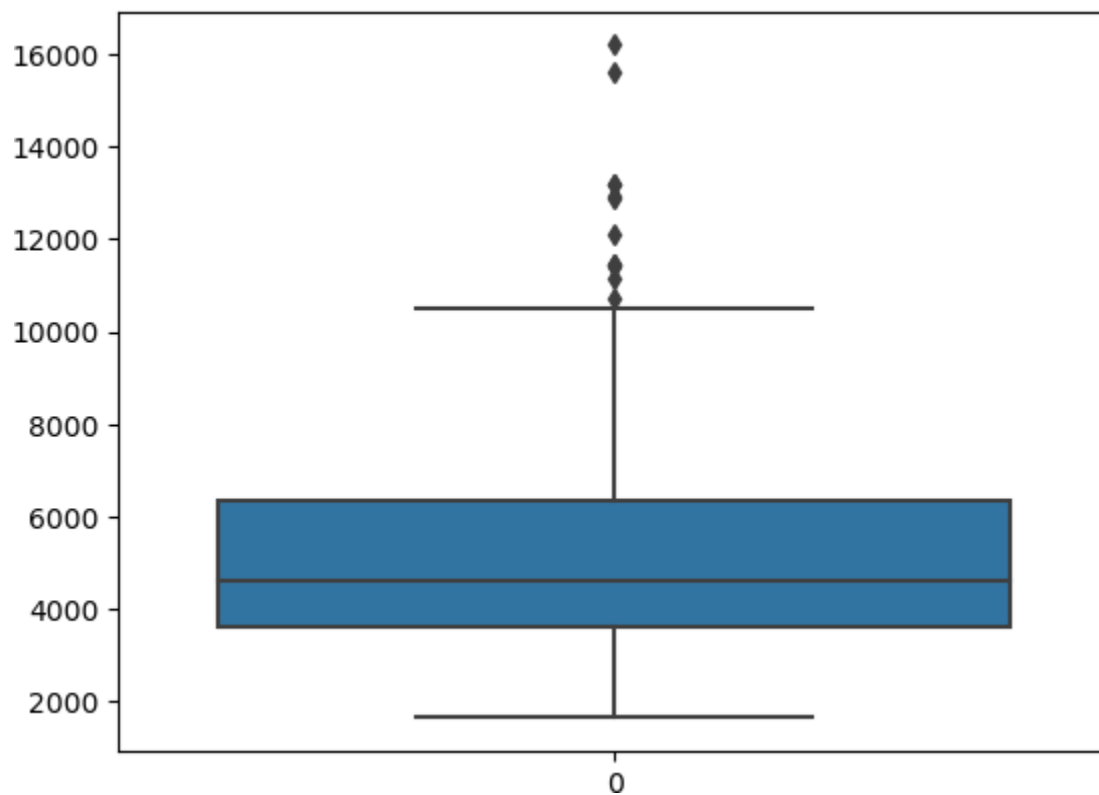
```
In [23]: <Axes: >
```

Out[23]:




```
In [24]: sns.boxplot(df['area'])
```

```
Out[24]: <Axes: >
```




```
In [25]: median_age = df['price'].median()
df["price"] = np.where(df["price"] > 12000000, median_age, df['price'])
sns.boxplot(df['price'])
```

```
Out[25]: <Axes: >
```



```
In [26]: median_area = df['area'].median()
df["area"] = np.where(df["area"] > 14000, median_area, df['area'])
sns.boxplot(df['area'])
```

Out[26]: <Axes: >



Check for Categorical columns and perform encoding.

```
In [27]: from sklearn.preprocessing import OneHotEncoder
```

```
In [28]: encoding = pd.get_dummies(df, columns = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'furnishingstatus'])
```

```
In [29]: encoding.head()
```

```
Out[29]:
```

	price	area	bedrooms	bathrooms	stories	parking	mainroad_no	mainroad_yes	guestroom_no	guestroom_yes
0	4340000.0	7420.0	4	2	3	2	0	1	1	0
1	4340000.0	8960.0	4	4	4	3	0	1	1	0
2	4340000.0	9960.0	3	2	2	2	0	1	1	0
3	4340000.0	7500.0	4	2	2	3	0	1	1	0
4	11410000.0	7420.0	4	1	2	2	0	1	0	1

Split the data into dependent and independent variables

```
In [30]: df.columns
```

```
Out[30]: Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking', 'furnishingstatus'], dtype='object')
```

```
In [65]: # independent variables
X = encoding.drop(['price'], axis = 1)
X.head()
```

```
Out[65]:
```

	area	bedrooms	bathrooms	stories	parking	mainroad_no	mainroad_yes	guestroom_no	guestroom_yes	b
0	7420.0	4	2	3	2	0	1	1	0	
1	8960.0	4	4	4	3	0	1	1	0	
2	9960.0	3	2	2	2	0	1	1	0	
3	7500.0	4	2	2	3	0	1	1	0	
4	7420.0	4	1	2	2	0	1	0	1	

```
In [66]: # dependent variables
y = df[['price']]
y.head()
```

```
Out[66]:
```

	price
0	4340000.0
1	4340000.0
2	4340000.0
3	4340000.0
4	11410000.0

Scaling the independent variables

```
In [67]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_std = scaler.fit_transform(X)
```

```
In [68]: x_std
```

```
Out[68]: array([[ 1.11756482,  1.40341936,  1.42181174, ...,  1.70084013,
                -0.84488844, -0.6964292 ],
                [ 1.8623093 ,  1.40341936,  5.40580863, ...,  1.70084013,
                -0.84488844, -0.6964292 ],
                [ 2.34590961,  0.04727831,  1.42181174, ..., -0.58794474,
                 1.18358821, -0.6964292 ],
                ...,
                [-0.72011635, -1.30886273, -0.57018671, ..., -0.58794474,
                -0.84488844,  1.43589615],
                [-1.06347257,  0.04727831, -0.57018671, ...,  1.70084013,
                -0.84488844, -0.6964292 ],
                [-0.60888828,  0.04727831, -0.57018671, ..., -0.58794474,
                -0.84488844,  1.43589615]])
```

Split the data into training and testing

```
In [69]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=0)
```

Build the Model

```
In [70]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [71]: lr = LinearRegression()
```

```
In [ ]:
```

Train the Model

```
In [72]: lr.fit( X_train, y_train )
```

```
Out[72]: ▼ LinearRegression
LinearRegression()
```

```
In [73]: print("Value of the coefficients: \n", lr.coef_)
print("-----")
print("Value of the intercept: \n", lr.intercept_)
```

```
Value of the coefficients:
[[ 2.60781675e+02  9.30932038e+04  8.20605321e+05  3.96961106e+05
   1.00796216e+05 -3.10469714e+05  3.10469714e+05 -3.32132424e+05
   3.32132424e+05 -1.72635846e+05  1.72635846e+05 -7.33905647e+05]
```

```
7.33905647e+05 -4.83024979e+05 4.83024979e+05 1.04827468e+05
1.38634062e+05 -2.43461530e+05]]
```

```
-----
Value of the intercept:
[2156139.9017023]
```

Test the Model

```
In [74]: Y_pred = lr.predict(X_test)
```

Measure the performance using Metrics.

```
In [81]: from sklearn.metrics import mean_squared_error
        from sklearn.metrics import r2_score
```

```
In [82]: print("MSE",mean_squared_error(y_test,Y_pred))
        print("-----")
        print("RMSE",np.sqrt(mean_squared_error(y_test,Y_pred)))
        print("-----")
        print("R-Square", r2_score(y_test,Y_pred))
```

```
MSE 1326521791171.129
```

```
-----
RMSE 1151747.2774750236
```

```
-----
R-Square 0.5718914765881087
```

```
In [ ]:
```